


Regression

The SleepHealthLifestyle data set found on Kaggle data website was analyzed in this exploratory data analysis.

Using Linear Regression



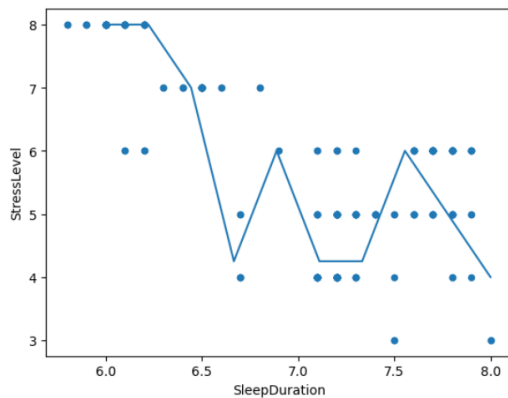
The relationship between duration of sleep and stress levels was studied.

Above, a linear regression model that predicts a person's Stress Level based on their Duration of Sleep (Number of Hours of Sleep Received) was built, along with a scatter plot and fitted line as a visual representation.

As shown in the model above, the fitted line is sloping downwards. The stress levels tend to be higher when someone has had less hours of sleep, while the stress levels tend to be lower when someone has had more hours of sleep. Thus, this visual shows that there is a relationship between Sleep Duration and Stress Levels.

Using K Nearest Neighbors Regression Function

```
[31]: <Axes: xlabel='SleepDuration', ylabel='StressLevel'>
```



Above, a k nearest neighbors regression model with 4 neighbors was built to predict Stress Levels based on Sleep Duration in the span between 6 – 8 hours.

The visual demonstrates the relationship between Sleep Duration and Stress Levels. However, this visual is different from the linear regression model as this visual displays the resulting Stress Levels of individuals in a piecewise manner because the goal is to predict the Stress Levels using the idea of nearest neighbors.

For example, the Stress Levels for those whose sleep duration is between 6.0 – 6.5 hours is 8 as those data points in the similar sleep duration period have the same 4 nearest neighbors when predicting stress levels. Another example is that those with a sleep duration in the range between 7.0 – 7.5 hours all have a predicted Stress Level of approximately 4.5, as demonstrated by the fitted line. Data points near the 7.0 – 7.5 hour Sleep Duration range have the same nearest neighbors, which is why the Stress Levels are predicted to be approximately 4.5 for all those data points, making the stress level prediction constant for this range of sleep duration periods.

Comparing the Results of Stress Levels from SleepDuration: Linear Regression Vs. K Nearest Neighbors:

The K Nearest Neighbors model to predict Stress Levels could potentially be a more accurate depiction, compared to the Linear Regression Model.

The linear Regression Model does a great job at conveying the general idea that the more hours a person spend sleeping, the lower their stress levels will be. The line shows a clear negative correlation between Sleep Duration and Stress Levels.

However, the fitted line in the Linear Regression model fails to accurately show that data points in a similar SleepDuration range have similar StressLevels. While the individual data points in the linear model all lie next to each other, the fitted line does not group them together appropriately for analysis. The K Nearest Neighbors model is more effective for

this as it groups the data in a piecewise manner based on similarity of Stress Levels and Sleep Duration using K Nearest Neighbors. From the K Nearest Neighbors model, conveys that those who sleep a certain amount of hours will have a very similar Stress Level, which is how this model is more accurate.

RMSE & MAE SleepDuration & StressLevel

The data was fitted by creating a pipeline that included KNeighborsRegressor with `n_neighbors = 4`. The Skicit package was then used to calculate the MSE and the RMSE.

As a result, RMSE had a value of 0.64.

The RMSE of 0.64 tells that the K Neighbors model prediction of Stress Levels based on Duration of Sleep is off by an average of 0.64. This implies that the using K Nearest Neighbors model to fit the data on Stress Levels may not be the most accurate.

As shown by the Mean Absolute Error, the average variance between Predicted Stress Levels and Actual Stress Levels is 0.43. The MAE of 0.43 is not terrible, however the K Neighbors model definitely does not predict Stress Levels perfectly.

Section 5.5 K Fold Cross Validation

Now, 2 variables were used to predict Stress Level.

The SleepDuration and SleepQuality variables from the data set were used to make predictions on StressLevel.

The train data set comprised of the first 200 rows of the original data set. The data was then randomized, and 50% of the random data went into train, and 50% of the random data went into val. The train data set was used to predict val. A pipeline was then made using K Nearest Neighbors value of 4, and the pipeline was fitted using the train data. The *predict* function on the pipeline was used to predict the y values on the validation data sets.

The Root Mean Square error was then computed, and resulted in a value of 0.61. This means that the average difference between the actual StressLevel and the predicted

StressLevel is around 0.61. 0.61 is a relatively high RMSE, so this model does not necessarily predict the StressLevel precisely from SleepDuration and SleepDuration.

Basically, the idea of cross validation is to use the first half of the training dataset to predict the 2nd half of the data, or use the 2nd half of the training dataset to predict the first half of the training dataset.

In the photograph above, the train data (1st half) was used to predict the val data (2nd half) and the val data (2nd half) was used to predict the train data (1st half). The RMSE for both was obtained. Their RMSE is not the same, as an RMSE of 0.61 was obtained when using train to predict val, and an RMSE of 0.525 when using val to predict train. This means that using val to predict train for StressLevel is more accurate than using train to predict val for the StressLevel Variable.

Cross Validation Section 5.5

Apparently, using half of the data in cross validation to predict Stress Levels does not necessarily give accurate results, so K Fold Cross Validation Method was used alternatively to try and predict StressLevel.

4 folds were made, where 3 subsamples are used as training data, and 1 subsample represents val, which is the data that will be predicted. The RMSE was computed after using K Fold Cross Validation to predict StressLevel, which was 0.67. Surprisingly, the RMSE was way higher compared to using cross validation.

However, the CV changed from CV = 4 to CV = 100, and as a result, the RMSE had significantly lowered to 0.40 when using 100 folds instead of just 4 folds. This RMSE was significantly lower than the RMSE found for cross validation.


In conclusion, at least 100 Folds are needed for the model to accurately predict StressLevels based on SleepQuality and SleepDuration.

Model Selection and Hyperparameter Tuning

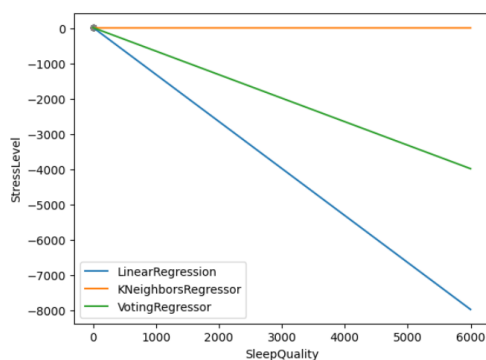
Hyperparameter Tuning was now used to determine the number of nearest neighbors that will minimize the Mean Square Error in order to accurately predict StressLevel from SleepQuality and SleepDuration.

A K Fold Cross Validation model was fitted with 100 folds, as specified by $cv = 100$. As a result, the number of neighbors needed to get the most accurate prediction of StressLevel is 16 neighbors, as the MSE is minimized at 16 neighbors.

Lastly, a loadings plot with 16 neighbors was made.



The first loading plot above represents the relationship between SleepDuration and StressLevel.



[]:

The second loading plot represents the relationship between sleep quality and stress level.

The loading plot displays that as Sleep Quality and Sleep Duration increases, Stress Levels tend to decrease. These are shown from Linear Regression and Voting Regression methods. However, the K Neighbors Regressor Line is completely straight, which does not seem to be an accurate depiction between the variables.

Classification

6.1 K Nearest Neighbors for Classification

For classification, the BMI category that a person would fall under would be predicted based on their how long they sleep (SleepDuration) and the quality of their sleep (SleepQuality).

A 5 neighbors model was fitted with the X variables being SleepDuration and SleepQuality, and Y variable being “BMI Category”.

First, the probabilities for each BMI category were predicted using sleep Duration of 5.9 Hours and Sleep Quality of 6. The results show that 80% of these people were overweight, and 20% of these people were normal in weight. According to value counts, the majority of people with these sleep statistics were overweight.

```
>]: BMICategory
Overweight    3
Normal        2
```

When using 10 nearest neighbors to predict BMICategory from SleepDuration and SleepQuality, 90% of the people were overweight and only 10% of these people were normal in weight.

When utilizing 10 neighbors for prediction, Value Counts showed that 7 of the people were overweight and 3 of the people were of normal weight.

```
[9]: BMIcategory
      Overweight    7
      Normal        3
```

Using 100 Nearest Neighbors:

```
] BMIcategory
      Overweight    66
      Normal        33
      Obese          1
```

Sleep Duration of 5.9 hours and Sleep Quality of 6.

A model was fitted to predict the BMI categories that people fall under when they have a sleep duration of 5.9 hours and a Sleep Quality rating of 6. The findings show that 66% of these people fall under the Obese BMI category.

```
[36]: BMIcategory
      Normal        58
      Overweight    33
      Normal Weight  7
      Obese          2
```

Sleep duration of 8.1 hours with SleepDuration and SleepQuality of 9

A model was fitted to predict the BMI categories that people fall under when they have a sleep duration of 8.1 hours and a sleep quality of 9 hours. 100 nearest neighbors were used, as more data was obtained for each of the categories. The findings show that 58% of the people were normal in weight, and 2% of these people were obese. This shows that those who get more sleep are generally normal in weight.

The comparisons when using 100 nearest neighbors show that those who sleep for more hours and have better quality sleep tend to fall under the normal weight BMI category. However, those who sleep for less hours and have poor quality of sleep tend to fall under the Obese BMI category.

Earlier, predictions were done using 5 nearest neighbors and 10 nearest neighbors. However, the analysis using 100 nearest neighbors is a more accurate depiction of how sleep quality, sleep duration, and the BMI category that a person would fall under is related.

6.2 and 6.3 Precision, Recall, and F1 Score

Now, precision and recall were calculated for the three BMICategories: Normal, Overweight, and Obese.

The BMI Category of Normal has a **precision score** of 0.799 and a **recall score** of 1.0. The **precision score** of 0.799 means that approximately 80% of the observations that were predicted to fall under the Normal BMI category fell under the Normal BMI category correctly. The recall score of 1.0 means that 100% of the observations that actually fell under the Normal BMI category were predicted to be in the Normal BMI Category.

The BMI Category of Obese has a **precision score** of 1.0 and a **recall score** of 0.4. The **precision score** of 1.0 means that approximately 100% of the observations that were predicted to fall under the Obese BMI category fell under the Obese BMI category correctly. The recall score of 0.4 means that 40% of the observations that actually fell under the Obese BMI category were predicted to be in the Obese BMI Category.

The BMI Category of Overweight has a **precision score** of 0.9 and a **recall score** of 0.736. The **precision score** of 0.9 means that approximately 90% of the observations that were predicted to fall under the Overweight BMI category fell under the Overweight BMI category correctly. The recall score of 0.736 means that 73.6% of the observations that actually fell under the Overweight BMI category were predicted to be in the Overweight BMI Category.

The F1 scores were computed for each of the BMI categories to get a general idea of their accuracy.

As a result, f1 scores were collected for each of the BMI categories. The Normal BMI category has an f1 score of 0.88, and the Overweight BMI category has an f1 score of 0.81.

These scores imply that the predictions for the Normal and Overweight BMI categories are relatively accurate, however, either the precision or recall score for one of these should be slightly higher to ensure more accuracy.

Unfortunately, the f1 score for the Obese category is around 57%, which is pretty low in accuracy compared to Normal and Overweight BMI. This means that either the precision or recall score is significantly low. In order to increase the f1 score, the precision or recall score must be increased.

Hyperparameter tuning was then used to obtain a K value that will lead to the best f1 Score.

As a result, using 3 neighbors will give the highest f1 Score. This means that using 3 neighbors will most accurately predict the BMI category that a person falls under based on the variables SleepQuality, SleepDuration, and StressLevel.